

Modernizing C++ Legacy Systems with Rust: Strategies for Safe Interoperability and Performance Evaluation in Rust-Based Boost Components

Filipe Daniel Ferreira

**Dissertation Research Plan of the Master's degree in Critical Computing
Systems Engineering**

Supervisor: Luis Miguel Pinho

Porto, February 1, 2026

Statement of Integrity

I hereby declare having conducted this academic work with integrity. I have not plagiarised or applied any form of undue use of information or falsification of results along the process leading to its elaboration.

I declare that the work presented in this document is original and my own, and has not previously been used for any other purpose.

I further declare that I have fully followed the Code of Good Practices and Conduct of the Polytechnic Institute of Porto.

ISEP, Porto, February 1, 2026

Abstract

The growing complexity of safety-critical software systems, particularly in the automotive domain, has intensified concerns regarding memory safety, concurrency correctness and long-term maintainability of large C++ codebases. Although C++ remains the dominant language in embedded and real-time systems due to its performance and mature ecosystem, its reliance on manual memory management continues to be a major source of vulnerabilities. Rust has emerged as a promising alternative, offering strong compile-time guarantees for memory safety and data-race freedom without introducing garbage collection. However, the adoption of Rust in industrial environments is hindered by the scale of existing C++ systems and their deep dependence on template-heavy frameworks such as the Boost C++ Libraries.

This dissertation investigates strategies for incrementally modernizing legacy C++ systems by integrating Rust in a safe and performant manner. Focusing on Boost-based components commonly used in automotive software architectures, it proposes a systematic methodology for selecting migration candidates, designing bidirectional interoperability interfaces, and re-implementing critical functionality in Rust. The work explores Foreign Funtional Interfaces (FFI) patterns, evaluates existing interoperability tools, and addresses the architectural challenges arising from the mismatch between C++ templates and Rust's ownership model.

A Rust-based reimplementation of selected Boost components is developed and integrated into a hybrid C++/Rust system. Performance, memory usage, and interoperability overhead are experimentally evaluated to quantify the benefits and costs of this approach. The results demonstrate that incremental migration can significantly improve memory safety while preserving determinism and performance, providing practical guidance for the adoption of Rust in high-assurance and safety-critical software systems.

Keywords: Boost, C++, Rust, interoperability, migration, memory safety

Resumo

A crescente complexidade dos sistemas de software crítico e seguro, particularmente no domínio automóvel, intensificou as preocupações relativas à segurança da memória, à correção da concorrência e à manutenção a longo prazo de grandes bases de código C++. Embora o C++ continue a ser a linguagem dominante em sistemas embebidos e de tempo real devido ao seu desempenho e ecossistema maduro, a sua dependência da gestão manual da memória continua a ser uma importante fonte de vulnerabilidades. O Rust surgiu como uma alternativa promissora, oferecendo fortes garantias em tempo de compilação para a segurança da memória e a ausência de conflitos de dados, sem introduzir a recolha de lixo. No entanto, a adoção do Rust em ambientes industriais é dificultada pela escala dos sistemas C++ existentes e sua profunda dependência de frameworks fortemente baseadas em modelos, como as bibliotecas Boost C++.

Esta dissertação investiga estratégias para modernizar incrementalmente sistemas C++ antigos, integrando Rust de maneira segura e com bom desempenho. Com foco em componentes baseados em Boost comumente usados em arquiteturas de software automotivo, ela propõe uma metodologia sistemática para selecionar candidatos à migração, projetar interfaces de interoperabilidade bidirecionais e reimplementar funcionalidades críticas em Rust. O trabalho explora padrões de FFI, avalia ferramentas de interoperabilidade existentes e aborda os desafios arquitetônicos decorrentes da incompatibilidade entre modelos C++ e o modelo de propriedade do Rust.

Uma reimplementação baseada em Rust de componentes Boost selecionados é desenvolvida e integrada a um sistema híbrido C++/Rust. O desempenho, o uso de memória e a sobrecarga de interoperabilidade são avaliados experimentalmente para quantificar os benefícios e custos dessa abordagem. Os resultados demonstram que a migração incremental pode melhorar significativamente a segurança da memória, preservando o determinismo e o desempenho, fornecendo orientações práticas para a adoção do Rust em sistemas de software de alta garantia e segurança crítica.

Contents

1	Introduction	1
1.1	High-Level Overview and Motivation	1
1.2	Problem Statement	2
1.3	Research Questions	3
1.4	Research Objectives	3
1.5	Research Contributions	4
1.6	Dissertation Plan Structure	4
1.7	Privacy and ethics	4
1.8	Use of AI Tools	5
2	State of the art	7
2.1	Rust	7
2.1.1	Overview	7
	Ownership and Memory Safety	8
	Tooling and ecosystem	9
	Adoption and Industrial Applications	9
2.2	C++ and the Boost Library	10
2.2.1	C++	10
2.2.2	Boost C++ Libraries	12
2.2.3	The Boost Paradox: Barrier and Bridge	12
2.3	Interoperability between Rust and C++	16
2.3.1	Undefined Behavior and Vulnerability Mitigation	16
2.3.2	The Safety Air Gap in Hybrid Architectures	16
2.3.3	Systematic Porting vs. the Rewrite Trap	17
2.3.4	Foreign Function Interface (FFI)	17
2.3.5	Interoperability tools	18
2.3.6	Boost migration for Rust	18
2.4	Case studies	19
2.4.1	Concurrency Safety: Migrating Locking primitives	19
2.4.2	“Just Use Rust”: A Best-Case Historical Study of Open Source Vulnerabilities in C	20
3	Methodology and Solution Proposal	23
3.1	Component Selection Criteria	23
3.1.1	Automotive-Oriented Selection Rationale	23
3.1.2	Deprioritized Components	24
3.1.3	Safety-Critical Constraints and ISO 26262 Alignment	24
3.1.4	Boost Libraries Relevant to Eclipse Safety-Critical Open Research Environment (SCORE)	25
3.1.5	Comparative Analysis as a Migration Enabler	25
3.1.6	From Library Gaps to New Rust Automotive Crates	26

3.1.7 Automotive Safety Integrity Levels (ASIL)-Oriented Migration Strategy . .	27
3.1.8 Migration Workflow Overview	27
4 Development Plan	29
4.1 Objectives and Scope	29
4.2 Development Phases Overview	29
4.3 System Analysis and Component Selection	30
4.4 Architectural Refactoring and Interface Design	30
4.5 Rust Reimplementation	31
4.6 Bidirectional FFI Integration	31
4.7 Validation and Safety Analysis	31
4.8 Consolidation and Documentation	32
4.9 Development Timeline	32
Bibliography	33

List of Figures

2.1	Concrat	20
2.2	rustpreventions	21
3.1	Boost-to-Rust Migration Workflow	28

List of Tables

2.1	Rust ecosystem crates that replace common Boost C++ Libraries.	19
2.2	Top 20 Most Common Vulnerabilities and Rust Prevention Status (Meneely et al. 2025)	22
3.1	Boost libraries relevant to Eclipse SCORE-style automotive systems and Rust equivalents	26
4.1	Gantt-style development plan for the Boost-to-Rust migration	32

List of Abbreviations

List of Acronyms

ASIL	Automotive Safety Integrity Levels.
CWE	Common Weakness Enumeration.
ECU	Electronic Control Units.
FFI	Foreign Funtional Interfaces.
QM	Quality Management.
RAII	Resource Acquisition Is Initialization.
SCORE	Safety-Critical Open Research Environment.
SDV	Software Defined Vehicles.
STD	Standard C++ library.
TMP	Template Metaprogramming.
WCET	Worst Case Execution Time.

Chapter 1

Introduction

1.1 High-Level Overview and Motivation

The automotive industry is performing a profound transformation driven by the rise of the Software Defined Vehicles (SDV), autonomous driving technologies and increasingly interconnected Electronic Control Units (ECU) (S&P Global Mobility 2025). With the growth of these systems in complexity and responsibility, the reliability of high-assurance software becomes critical. The systems must operate under strict real-time, reliability and safety requirements regulated by standards as ISO 26262 (Deloitte 2024). In this environment, memory safety has a special importance, as software faults can directly translate into physical hazards (Bristol 2025).

Historically, C++ has been the dominant language (paired with C) in embedded and automotive development due to its performance, deterministic resource use and mature ecosystem (*ISO 26262-6:2018 Road vehicles — Functional safety — Part 6: Product development at the software level* 2018). However, its dependence on manual memory management continues to be a major source of vulnerabilities such as: buffer overflows, dangling pointers, race conditions and undefined behaviors (The MISRA Consortium 2023). Such issues are incompatible with the needs of safety critical cyber-physical systems, where a minor software defect can lead to serious consequences (AUTOSAR Partnership 2023).

Rust has rapidly emerged as a plausible successor for safety-critical software. Enforcing memory safety, ownership and data race free concurrency at compile time (*Rust Documentation* 2025). Rust promises to drastically reduce classes of defects endemic to C/C++ systems (Blandy, Orendorff, and Tindall 2021). Rust achieves this without introducing garbage collection, retaining the predictability required for real-time environments (*Rust Documentation* 2025).

Despite the advantages Rust presents, the transition from C++ to Rust in large industrial ecosystems remains complex. Automotive software contains millions of lines of validated legacy C++ code. This code is often built at top of complex frameworks and libraries as Boost C++ libraries (Tolnay 2025). Rewriting those systems from scratch is infeasible. To avoid it, the companies usually pursue incremental and hybrid migration. This migration requires robust interoperability between Rust and C++ (Google Android Team 2024).

This dissertation is motivated by the need to address that gap: enabling safe, practical and measurable integration of Rust into deeply entrenched C++ based automotive architectures, specially real-time related architectures.

Rewriting millions of lines of validated, legacy C++ code is operationally infeasible, necessitating a hybrid integration strategy. This transition is particularly problematic for automotive subsystems (such as sensor fusion, V2X communication and path planning) that are deeply coupled with the C++ ecosystem and reliant on complex, template-heavy frameworks like the Boost C++ Libraries.

1.2 Problem Statement

The central results from the tension between two different realities:

1. **C++ is deeply entrenched in automotive systems**, specially through heavily templated, performance-critical frameworks like Boost.
2. **Rust provides strong memory safety guarantees**, but its adoption is hindered by the difficulty of interoperating with complex C++ abstractions creating a challenge for a transition between C++ and Rust.

Currently, the interoperability approaches rely mainly on C-style Foreign Funtional Interfaces (FFI) . However, it is not possible to natively represent C++ templates, ownership semantics, RAII lifecycles or advanced Boost abstractions with C FFI (Google 2021).

As result of this limitation:

- The integration often requires unsafe code, reducing Rust's safety benefits.
- Interfacing with templaty heavily libraries, as Boost.Geometry and Boost.Asio, introduces high engineering overhead.
- Crossing the Rust and C++ boundary imposes a performance cost that is poorly characterized for high-assurance systems.
- No established methodology exists for the systematic and incremental migration of C++ Boost components to Rust while maintaining compatibility with real-time automotive subsystems.

The problem is not simply porting code from C++ to Rust, but the absence of validated techniques, architectural patterns, and tooling to bridge the impedance mismatch between C++ templates and Rust's ownership model, without breaking determinism, performance or safety requirements.

1.3 Research Questions

To address this gap, this dissertation is guided by the following research questions:

1. How can Rust be integrated into existing C++/Boost based automotive systems in a way that preserves and improves safety, determinism and performance?
2. What are the technical criteria required to identify which Boost library components are the most viable candidates for a migration to Rust within an automotive context?
3. Which Boost components represent suitable candidates for partial or complete migration to Rust, considering complexity, memory safety risk and template dependency?
4. What architectural patterns and FFI strategies enable safe and efficient bidirectional communication between Rust and C++?
5. What measurable performance and memory safety benefits does a Rust based reimplementation provide when compared with its Boost counterpart?
6. What are the interoperability costs ("FFI tax") and how do they impact hybrid real-time systems?

1.4 Research Objectives

The primary objective of this dissertation is to design, implement and evaluate a methodology for the incremental migration of Boost C++ Boost components into Rust, ensuring safe and efficient interoperability within hybrid systems. This involves re-architecting critical components to leverage Rust's memory safety guarantees and evaluating resulting architectural trade-offs.

The specific objectives are

- Analyze the current interoperability mechanisms between Rust and C++ (e.g., CXX, Bindgen, autocxx, cbindgen).
- Identify and define criteria for selecting Boost components suitable for migration (eg. memory-safety relevance, template complexity, use in automotive systems).
- Re-architect and re-implement a selected Boost component in Rust while preserving semantics and meeting strict ownership and borrowing guarantees.
- Design bidirectional Rust and C++ interfaces enabling both C++ to call Rust and Rust to call C++ components.
- Measure runtime performance, memory footprint and the overhead of FFI interactions.
- Evaluate the overall feasibility, benefits and limitations of incremental C++ to Rust migration in high-assurance environments.

1.5 Research Contributions

The contributions of this dissertation are:

- A comprehensive overview of Rust, C++ and Rust interoperability and the Boost library ecosystem.
- A systematic analysis of existing interoperability mechanisms and their suitability for template-heavy automotive codebases.
- A formalized methodology for selecting and migrating Boost components to Rust.
- A Rust-native reimplementation of selected Boost component with full memory safety guarantees.
- A robust bidirectional FFI integration layer
- A Performance evaluation comparing Rust and C++ implementations and quantifying overhead

1.6 Dissertation Plan Structure

The remainder of this dissertation plan is organized as follows:

- **Chapter 2 - State of the Art:**
Presents the technological background and related work. Includes Rust fundamentals, C++ interoperability mechanisms, Boost library architecture and prior migration initiatives.
- **Chapter 3 - Methodology and Solution proposal:**
Defines the criteria for selecting Boost components and describes the proposed migration methodology and architectural designed.
- **Chapter 4 - Development Plan:**
Details the Rust reimplementation, interface design and bidirectional FFI integration.

1.7 Privacy and ethics

In the context of this dissertation plan, which is focused in modernizing C++ based systems using Rust, ethical considerations were primarily centered on scientific integrity, software safety and intellectual property compliance.

The research relied on public open-source codebases, as Boost libraries. No personal data, user-identifiable information (PII), or private operational logs were collected, processed, or stored. Consequently, the study falls outside the scope of GDPR requirements regarding human subjects.

1.8 Use of AI Tools

For the preparation of this dissertation plan and supporting documentation, AI tools (specifically large language models) were used strictly for the following purposes:

- **Document writing evaluation and improvement**, including linguistic refinement, clarity enhancement, and reorganization of text originally written by the author.
- **Assistance in identifying potential sources and directions for literature review**, such as suggesting relevant topics, keywords, and areas of research to investigate. All cited references included in the dissertation were individually verified, read, and selected by the author.
- Structural feedback and general writing suggestions for non-technical sections.
- Code example generation to support the investigation of the state of the art and also to facilitate the understanding.

The tools used were Perplexity (latest version), for deep search specialty for community based information, chat GPT (version GPT5.1) and Gemini (version 3), for document writing evaluation and improvement.

No AI-generated technical explanations, conceptual developments, design decisions, code, or analysis results were incorporated directly into the dissertation. All technical content is exclusively authored by the candidate.

The author takes full responsibility for the correctness, originality, and academic integrity of the work presented. If additional AI tools are used during the subsequent stages of the dissertation, this section will be updated accordingly to specify which tools were used, for what purpose, and in which components of the document. If no further AI usage occurs, this will also be explicitly stated.

Chapter 2

State of the art

The integration of Rust into C++ environments, particularly Boost libraries, marks a significant evolution in software development. Rust's core strengths are memory safety and concurrency without garbage collection, addressing longstanding vulnerabilities in C++ systems. Recent studies on microcontroller applications show that Rust can have similar speed to C with lower energy use than interpreted languages. While performance evaluations on platforms, such as Raspberry Pi, indicates that Rust FFT implementations can match or exceed C in energy efficiency (Dashdamirli 2025; Rooney and Matthews 2023).

These characteristics are dominant in systems with strict deterministic requirements, such as automotive Electronic Control Units (ECUs), high-frequency trading platforms, and real-time robotics. In these domains, the non-determinism introduced by garbage-collection "stop-the-world" pauses is not appropriate due to hard real-time latency constraints. Furthermore, resource-constrained embedded environments often cannot afford the runtime overhead associated with managed languages. Rust solves this problem through an ownership model and a strict borrow checker enforced at compile time. By tracking the lifetime and scope of every value during compilation, Rust guarantees memory safety and data-race freedom without requiring a runtime garbage collector. (Okawa and Yamaguchi 2024) This ensures that the resulting binary maintains the predictable performance and low resource footprint of C++, while eliminating entire classes of bugs such as null pointer de-referencing and use-after-free errors.

2.1 Rust

2.1.1 Overview

Rust is a systems programming language focused on three goals: safety, speed, and concurrency (Klabnik, Nichols, and Developers 2017).

It is designed to ensure memory safety and high performance without relying on a garbage collector, present in languages as Python or Java.

To not rely on garbage collection and to avoid common memory issues found in languages like C and C++, such as null pointer dereferencing, double frees, and segmentation faults, Rust was designed to be memory safe while offering comparable performance (Khan 2023).

Ownership and Memory Safety

Rust distinguishes itself through a rigorous ownership model, which manages memory through a strict set of rules enforced at compile time by the "borrow checker", without a garbage collector. (Blandy, Orendorff, and Tindall 2021)

The garbage collector, used by Python and Java, is an automatic memory management process that helps programs to run efficiently. Rust does not use it for memory management, instead it determines exactly when memory should be freed based on variable scope. (Blandy, Orendorff, and Tindall 2021)

Within a program run, every value has a single "owner" and when the owner goes out of scope the value is dropped. Ownership can be transferred (moved) but not implicitly copied for complex types, preventing double-free errors. (Blandy, Orendorff, and Tindall 2021)

It is also possible to "borrow" values through references. The compiler ensures that references are always valid and a mutable reference cannot exist while other references exist, preventing data races. (Khan 2023)

Scope and ownership code example:

```
fn main() {
    // 1. SCOPE AND OWNERSHIP
    {
        let s1 = String::from("hello"); // s1 owns the string
        // s1 is valid here
    } // s1 goes out of scope here; memory is automatically freed (dropped)

    // 2. MOVE SEMANTICS (Transferring Ownership)
    let s2 = String::from("complex type");
    let s3 = s2; // Ownership moved from s2 to s3

    // println!("{}", s2); // Error! s2 is invalid. Prevents double-free errors.
    println!("s3 owns the data: {}", s3);

    // 3. BORROWING (References)
    let mut s4 = String::from("data");

    // Immutable borrow
    let r1 = &s4;
```

```

let r2 = &s4; // Multiple immutable references are allowed
println!("Read only: {} and {}", r1, r2);
// r1 and r2 scopes end here (last usage)

// Mutable borrow
let r3 = &mut s4;
r3.push_str(" race"); // Modify the borrowed value
// let r4 = &s4; // Error! Cannot have immutable ref while mutable ref exists
// This rule prevents data races.

println!("Modified data: {}", r3);
}

```

Tooling and ecosystem

An high quality tooling is a strong base of the Rust development environment. The integrated package manager **Cargo**, documentation generator **rustdoc** and built-in testing framework significantly enhance developer productivity. The Rust package registry, crates.io, is an important source of libraries for different applications.

Beside the growing usage of Rust, it still relies on interoperability with C/C++ ecosystem, specially for low-level applications (eg. hardware interactions, graphics, legacy protocols). Interfacing Rust with foreign code is possible using through the FFI. Usually, an FFI is functional and efficient but often requires carefull design of abstractions, specially when crossing language boundaries with different memory management models.

Adoption and Industrial Applications

Rust has been adopted by several domains, as:

- Web browsers and rendering engines, as mozilla firefox (Mozilla 2026).
- Cloud infrastructures, as AWS firecracker (Amazon Web Services 2026).
- Operating systems and kernel modules, as in Linux kernel (The Linux Kernel Organization 2026).
- Blockchain clients and cryptographic systems (Kuo 2024).
- High performance networking and embedded systems (The Rust Project Developers 2026).

2.2 C++ and the Boost Library

2.2.1 C++

C++ is currently the dominant programming language due to its extensive ecosystem. It has hardware specific libraries, mature compilers, profiling and debugging tools and frameworks for high-performance computing. Is also used for a large group of companies for many years (Stroustrup 2013). It leads to massive codebases, thousands of lines of code that make large rewrites infeasible (CppCon 2020).

Some common industries that usually use C++ are telecommunications, automotive, finance, aerospace, gaming, real-time control and large scale cloud systems (Stroustrup 2020).

The main challenges these systems have is their modernization. Modernizing them while ensuring the continuity, compatibility and performance is not easy. Due to this, languages that interoperate with C++ instead of replacing it are preferred.

The main strengths of C++ are:

- High performance (specially comparing with C) (Stroustrup 2020).
- Deterministic resource management through Resource Acquisition Is Initialization (RAII) (Stroustrup 2013).
- Support for control of hardware and memory (Stroustrup 2013).
- String abstraction mechanisms as templates, operator overloading and generic programming (Stroustrup 2013).
- Mature tools, compilers, profiling and debugging ecosystems (CppCon 2020).

The evolution of C++ ISO standard introduced modern features as smart pointers, concurrency primitives, modules and constant evaluation, for safety improvement and expressiveness (*ISO/IEC 14882:2020 Programming languages — C++* 2020).

C++ memory model and safety limitations

Beside all the advantages, C++ has memory and safety limitations. C++ is designed to give developers complete control over memory and resources lifetime. It leads to the following unsafe behaviors:

- Use after free and dangling pointers (*70% of Security Bugs Are Memory Safety Issues* 2020).
- Double frees (*70% of Security Bugs Are Memory Safety Issues* 2020).
- Buffer overflows and out of bound access (*70% of Security Bugs Are Memory Safety Issues* 2020).
- Iterator invalidation (Stroustrup 2013).
- Data races (Stroustrup 2013).

- Uninitialized memory reads (*70% of Security Bugs Are Memory Safety Issues* 2020).
- Reliance on undefined behaviors (Stroustrup 2013).

Industry data confirms that memory-unsafe languages (primarily C and C++) account for the majority of security vulnerabilities. Google reports that 70% of Chrome vulnerabilities stem from memory safety bugs (*Memory Safety in Android and Chrome* 2022), and Microsoft reports similar figures across Windows components (*70% of Security Bugs Are Memory Safety Issues* 2020). Although modern C++ attempts to mitigate issues through smart pointers and safer API design, legacy code and raw pointer usage remain widespread (CppCon 2020).

C++ concurrency and data races

C++ has a standardized model of memory since C++ 11. Through Standard C++ library (STD), C++ provides access to threads (for multi processing and for parallel execution) and atomics that are a powerful mechanism for managing concurrent access to shared data in multi-threaded applications. Also provides mutexes, locks and condition variables for safety memory access for multiple processes concurrently.

With all those features C++ provides a possibility of an high performance for parallel programs, but the correctness of C++ remains difficult. That is due C++ not being capable of statically prevent data races, C++ allows aliasing and mutation via multiple paths, exposes low level memory order operations and relies on the developer to ensure invariants (Stroustrup 2020).

At following example is possible to see two threads modifying the same resource, the global variable **counter**. The modification is concurrently leading to a undefined behavior.

```
#include <iostream>
#include <thread>
#include <vector>

int counter = 0; // Shared resource

void increment() {
    for (int i = 0; i < 1000000; ++i) {
        // Data race occurs here: read-modify-write is NOT atomic
        counter++;
    }
}

int main() {
    std::thread t1(increment);
    std::thread t2(increment);

    t1.join();
```

```

    t2.join();

    std::cout << "Final counter: " << counter << std::endl;
    // Expected: 2000000
    // Actual: Likely < 2000000 (e.g., 146377) due to race conditions [web:51]
    return 0;
}

```

The incremental of **counter**, **counter++**, is not atomic. Usually an incremental operation includes three steps: the load of the variable into a register, the value increment at the register and the value store back at the memory.

When this operation is not protected by locking (e.g. using threads) or atomic types (**std::atomic<int>**), both threads can load the same value from memory, increment it and store it, simultaneously. Since the initial value is the same, instead of expected increment by 2 it only will increment by 1.

2.2.2 Boost C++ Libraries

The Boost C++ Libraries are a comprehensive collection of peer-reviewed, portable, and open-source libraries that function as a testbed for the C++ Standard Library. Boost has heavily influenced modern C++ standards (C++11 through C++23), serving as the incubator for foundational components such as smart pointers, threading models, and regular expressions. Boost has a rigorous peer review process and commitment to portability made it a benchmark for high quality C++ library development. The capacity of early provision of production ready implementations of libraries, enabled developers to adopt it. (*Boost C++ Libraries Documentation* 2025)

2.2.3 The Boost Paradox: Barrier and Bridge

The Boost libraries have, historically, served as the proving ground for the C++ Standard Library. Boost provide high-quality, peer-reviewed, and largely header-only components that target portability and performance across platforms. (*Boost C++ Libraries Documentation* 2025) This ubiquity creates a unique dichotomy during migration, because the same abstractions that enabled modern C++ design patterns also amplify the complexity of introducing a Rust-based ownership model into existing codebases.

- **Technical mismatch:** Boost utilizes advanced C++ features such as template metaprogramming, complex inheritance hierarchies, SFINAE (Substitution Failure Is Not An Error) that is a principal in C++ where an invalid substitution of template parameters is not an error in itself, and expression templates across libraries as `Boost.SmartPtr`, `Boost.Spirit`, and `Boost.Asio` (*Boost C++ Libraries Documentation* 2025).

These constructs don't have direct equivalents in Rust's generic system and trait bounds, particularly when combined with partial specialization and ADL-dependent overload sets. As a consequence, standard automated binding tools (such as `bindgen` (The rust-bindgen Project Developers 2026)) are ill-suited to interface with Boost's header-only, template-heavy design: the absence of concrete, non-templated symbols at the ABI level often forces developers to introduce intermediate C-compatible facades.

These facades are typically handwritten wrappers that manually erase templates into plain C types or opaque handles, reintroducing the same classes of human error (lifetime mismatches, missing error propagation, accidental copying) that Rust aims to eliminate in the first place.

This code uses a template, `async_wait` from the library `boost::asio`. It is not a compiled symbol in the library, is generated only when used. Rust is not able to link to it directly.

```
// async_engine.hpp
#include <boost/asio.hpp>
#include <iostream>

class AsyncEngine {
public:
    AsyncEngine() : timer_(io_, boost::asio::chrono::milliseconds(500)) {}

    // PROBLEM: This method uses a C++ Template for the handler (Lambda).
    // Bindgen ignores this because it requires template instantiation.
    // It relies on SFINAE to check if the handler has the right signature.
    template <typename CompletionToken>
    void start_task(CompletionToken&& token) {
        timer_.async_wait(token);
        io_.run();
    }

private:
    boost::asio::io_context io_;
    boost::asio::steady_timer timer_;
};
```

- **Semantic opportunity:** Many Boost libraries were explicitly designed to encode disciplined resource management and concurrency patterns, that anticipate concerns later formalized by Rust. `Boost.SmartPtr` library encapsulates ownership and shared lifetime through types, such as `boost::shared_ptr` and `boost::scoped_ptr`.

The library `Boost.Thread` provides higher-level threading primitives and synchronization constructs beyond the C++03 baseline (*Boost C++ Libraries Documentation* 2025).

Higher-level facilities, such as `Boost.Program_options` and `Boost.Asio`, encourage clear separation of configuration, I/O and business logic. It aligns well with Rust's preference for explicit data flow and minimal global state.

The migration opportunity lies in systematically leveraging these “vocabulary types” as a pivot layer: raw pointers and ad hoc ownership conventions at subsystem boundaries are first refactored into Boost abstractions, which can then be mechanically mapped to Rust equivalents (for example, `boost::shared_ptr<T>` $\rightarrow$ `std::sync::Arc<T>`, `boost::optional<T>` $\rightarrow$ `Option<T>`, or `boost::variant` $\rightarrow$ a Rust enum).

Treating Boost as a semantic bridge rather than a legacy liability allows the migration process to progressively tighten ownership and error-handling contracts on the C++ side before introducing Rust, thereby reducing the impedance mismatch at the FFI boundary.

A simulation of a shared configuration system based, it is usually used in ROS/robotics systems.

```
#include <vector>
#include <string>
#include <iostream>
#include <boost/variant.hpp>
#include <boost/shared_ptr.hpp>
#include <boost/make_shared.hpp>

// 1. Define the sum type: It's EITHER an int OR a string
typedef boost::variant<int, std::string> ConfigValue;

// 2. Define the container: A shared list of these values
typedef std::vector<ConfigValue> ConfigList;
typedef boost::shared_ptr<ConfigList> SharedConfig;

// Visitor to print values (Required to safely unpack the variant)
struct Printer : public boost::static_visitor<void> {
    void operator()(int i) const { std::cout << "Speed: " << i << std::endl; }
    void operator()(const std::string& s) const { std::cout << "Mode: " << s << std::endl; }
};

int main() {
    // 3. Create shared data
    SharedConfig config = boost::make_shared<ConfigList>();
    config->push_back(100);           // Add int
```



```

        config->push_back("Autonomous");    // Add string

// 4. Share it (cheap copy, points to same data)
SharedConfig worker_ref = config;

// 5. Iterate and Visit (Type-safe access)
for (const auto& item : *worker_ref) {
    boost::apply_visitor(Printer(), item);
}

return 0;
}

// Rust equivalent using built-in equivalents (enum and match) instead of library templates
use std::sync::Arc;

// 1. Define the sum type (Enum)
enum ConfigValue {
    Speed(i32),
    Mode(String),
}

fn main() {
    // 2. Create shared data (Vec wrapped in Arc)
    // Note: Rust requires explicit variants (ConfigValue::Speed) unlike Boost
    let config = Arc::new(vec![
        ConfigValue::Speed(100),
        ConfigValue::Mode(String::from("Autonomous")),
    ]);

    // 3. Share it (cheap clone, points to same data)
    let worker_ref = config.clone();

    // 4. Iterate and Match (Type-safe access)
    for item in worker_ref.iter() {
        match item {
            ConfigValue::Speed(i) => println!("Speed: {}", i),
            ConfigValue::Mode(s) => println!("Mode: {}", s),
        }
    }
}

```

2.3 Interoperability between Rust and C++

The transition from C++ to Rust represents a necessary evolution for high-assurance software, driven by the critical need to eliminate memory safety vulnerabilities characteristic of C++'s manual memory management. However, for systems deeply rooted in the C++ ecosystem, particularly those dependent on the Boost C++ Libraries, migration is not merely a syntactic translation but a complex architectural reconciliation. The core challenge lies in bridging two distinct paradigms without sacrificing the stability of massive legacy investments.

2.3.1 Undefined Behavior and Vulnerability Mitigation

A primary driver for this migration is the prevalence of undefined behavior in legacy C/C++ codebases, which often manifests as security vulnerabilities. A best-case historical study of open-source vulnerabilities indicates that approximately 58.2% of historical vulnerabilities in prominent C projects would have been virtually impossible in safe Rust. (Meneely et al. 2025)

Code vulnerabilities are formally listed at Common Weakness Enumeration (CWE). CWE was community developed and list common software and hardware weakness types.

Some critical vulnerabilities detected were:

- **Use-after-free (CWE-416):** While Boost smart pointers manage object lifetimes, they cannot strictly enforce compile-time borrow checking, leaving potential for dangling references if raw pointers are extracted. Rust's borrow checker eliminates this entirely in safe code.
- **Buffer overflows (CWE-120/125):** C++ containers (including `std::vector` and `boost::container`) rely on developer discipline to avoid out-of-bounds access. In contrast, Rust enforces bounds checking by default, preventing memory corruption.
- **Data races (CWE-362):** Rust's `Send` and `Sync` traits ensure thread safety at the type level, preventing data races that are common in complex C++ multithreaded applications, even when using `boost::thread`.

However, automated translation tools often fail to fully capitalize on these safety guarantees. Transpilers like C2Rust may preserve the unsafe semantics of the original C code, necessitating further static analysis or refactoring to achieve true safety. (Hong 2023; Ling et al. 2022)

2.3.2 The Safety Air Gap in Hybrid Architectures

Hybrid architectures are structural design patterns that allow two languages to coexist in a single system, as systems with C++ and Rust. They share the same low-level control but have

different compilation models and safety guarantees. Due to it the architecture should focus on isolating the boundary to minimize friction.

Since big-bang rewrites, a software strategy where a system is rebuilt from scratch (rewriting a C++ system in Rust), are resource-prohibitive, organizations are forced into incremental migration.

This creates a prolonged hybrid state where Rust modules must coexist with legacy C++:

- **FFI vulnerabilities:** Rust's safety guarantees (ownership, borrow checking) strictly end at the Foreign Function Interface (FFI) boundary. Crossing this boundary destroys safety unless rigorously wrapped, creating a safety air gap. Dynamic analysis tools like Rustcheck have been proposed to detect memory safety issues specifically arising from unsafe Rust usage at these boundaries. (Xia, Wu, and Hua 2023)
- **Integration risks:** Passing complex Boost data structures across the FFI boundary introduces risks of unsafe pointer usage, exception mismatches, and memory leaks.
- **Concurrency conflicts:** Bridging Boost.Asio or Boost.Thread with Rust's concurrency model can lead to undefined behavior, specifically when distinct async runtimes or threading models attempt to interoperate without a unified strategy.

2.3.3 Systematic Porting vs. the Rewrite Trap

A lack of structured methodology often leads to the rewrite trap stalled development due to the sheer complexity of porting feature-rich systems.

- **Performance overhead:** FFI boundaries introduce latency (10–20% overhead in some cases) due to data marshalling. In latency-sensitive domains (HPC, HFT), an ad hoc porting strategy that places FFI boundaries in hot loops is unacceptable.
- **Missing methodology:** There is currently no standardized process for decomposing a Boost-heavy application into portable units. Developers need a systematic way to identify architectural seams where the cost of the FFI boundary is minimized and the safety value of Rust is maximized.

The central problem is not just the translation of code, but the architectural interoperability of two diverging memory models. The challenge is to define a strategy that stops treating Boost as legacy debt to be discarded and starts utilizing it as a stabilizing intermediate layer.

2.3.4 Foreign Function Interface (FFI)

A Foreign Function Interface (FFI) is a mechanism that enables a program written in one programming language (the host language) to call routines or make use of services written in another language (the guest language). FFIs are critical for system interoperability, allowing high-level languages to leverage legacy codebases, operating system APIs, and high-performance libraries often implemented in C or C++. (*C++ ABI Compatibility* 2025) In compiled languages

like Rust, FFI is often handled through static declarations. The `extern` block defines function signatures that adhere to the C application binary interface (ABI) and Rust uses attributes such as `#[link(name = "...")]` together with `#[repr(C)]` on data structures to ensure layout compatibility and correct symbol resolution. (*Rust Documentation* 2025)

2.3.5 Interoperability tools

There are two main tools for interoperability between C/C++ and Rust: **cxx** and **autocxx**.

Cxx provides a safe mechanism for calling C++ code from Rust and Rust code from C++, not subject to undefined behaviors and memory safety risks often introduced when using **bindgen** or **cbindgen** to generate C-style bindings. Unfortunately, beside this, C++ code remains 100% unsafe.

AutoCxx is a tool for calling C++ from Rust in a heavily automated form. It includes all fluent safety from Cxx generating interfaces automatically from C++ headers using a variant of **bindgen**.

2.3.6 Boost migration for Rust

Currently there is no official migration of Boost from C++ to Rust. Both ecosystems, C++ and Rust, are designed within different paradigms. Boost is described as a peer reviewed open source repository of C++ libraries. It is a massive, centralized collection that is used as testing ground for features eventually destined for STD.

Rust follows a different model. Instead of having a massive "Boost" crate, the community builds specialized, high-quality crates that are distributed through crates.io.

Boost, as previous reference, has an heavy use of Template Metaprogramming (TMP). It is not straightforward translation, the porting to Rust. That is due the fact of C++ and Rust handle generics differently. C++ templates are checked at compile time, if the code works. Rust generics are definition checked. Boost also relies on complex class hierarchies and inheritance. On other side, Rust uses composition, meaning a total architectural rewrite for Rust implementation of Boost. Boost also uses frequently smart pointers to manage complex object graphs and Rust uses ownership and borrowing that is a completely different approach. Those Rust rules are used to avoid the overhead of reference counting.

There is an entry at Crates.io for porting Boost libraries to Rust but there is no updates since 2021. At this moment has no source code (Zakaram 2023). On other side, there are some Rust native equivalents to Boost C++ at Crates.io that are safer and faster. At the following table is possible to check some equivalents (*crates.io: The Rust Community's Crate Registry* 2025).

Boost Library	Rust Equivalent	Notes
Boost.Asio	tokio, async-std	Async IO, networking, event loop
Boost.Beast	hyper, request	HTTP clientserver, WebSockets
Boost.Filesystem	std::fs, std::path	Native filesystem support in Rust stdlib
Boost.ProgramOptions	clap, argh, pico-args	Command-line parsing
Boost.Serialization	serde	General-purpose serialization framework
Boost.Regex	regex	Fast regular expression engine
Boost.Optional	Option<T>	Built-in language feature
Boost.Variant	enum, enum_dispatch	Tagged unions sum types
Boost.Any	dyn Any	Type-erased storage
Boost.Graph	petgraph	Graph algorithms and data structures
Boost.Geometry	geo, geo-types	Computational geometry
Boost.SmartPtr	Box, Rc, Arc	Memory management primitives
Boost.Thread	std::thread, crossbeam	Threading and concurrency
Boost.Lockfree	crossbeam, loom	Lock-free data structures
Boost.Coroutine	async/await (built-in), futures	Cooperative concurrency

Table 2.1: Rust ecosystem crates that replace common Boost C++ Libraries.

2.4 Case studies

After checking the IEEE paper library were not found specific case studies that directly migrate or benchmark Boost C++ libraries. The papers mainly focus on migrating C codebases or in specific FFI algorithms.

2.4.1 Concurrency Safety: Migrating Locking primitives

Legacy C software often uses `pthread_mutex_lock/pthread_mutex_unlock` to control concurrency. If C2Rust is used to migrate the C code directly to Rust, the resulting code still contains unsafe C style mutex operations, preventing Rust compiler from enforcing thread-safety guarantees (Immunant, Inc. 2025).

Was proposed, by Jaemin Hong, a solution to automatically migrate C programs to Rust with an high focus on concurrency. In order to solve C2Rust problems, was proposed a solution named Concrat. The solution Concrat consists on a aggregation of a C2Rust solution, a static analyzer and a Rust code transformer. The static analyzer performs the dataflow analysis of the code generated by the C2Rust to produce a summary about data and flow lock. The transformer replaces the lock API according to the previous summary.




Figure 2.1: Concrat workflow (Hong 2023)

The tool achieved valid compilable Rust code in 29 of 46 cases collected from the github, proving that ownership-based concurrency can be automated even for legacy codebases (Hong 2023). From the 46 cases, 5 of them were also compilable but with a need of a manual fixes. The remain 12 cases were considered failures classified into three categories: functions conditionally acquiring locks, function pointers, and pointers to locks (Hong 2023).

Concrat focuses on C rather than C++, although its findings are applicable at C++ and also applicable to Boost C++ migration. Libraries as Boost.Thread and Boost.Asio have a strong dependency on manual synchronization, raw pointers and complex ownership patterns. The study has demonstrated that a simple translation of that code to Rust, relying on C2Rust or C++ bindings, would preserve the safety risks unless concurrency semantics are explicitly reconstructed. Concrat's static analysis offers a practical blueprint for the reconstruction providing a lock analysis.

This case study illustrates how automated, semantics-aware transformation can bridge the gap between C++'s flexible but unsafe concurrency model and Rust's strict, verifiable approach, enabling incremental and safer integration of Rust into existing C++ ecosystems.

2.4.2 “Just Use Rust”: A Best-Case Historical Study of Open Source Vulnerabilities in C

In this paper the authors provide a historical security analysis of C language projects, based on the most common vulnerabilities of C and provide answers Rust based for those vulnerabilities. Those vulnerabilities were divide into two groups: the vulnerabilities Rust could prevent and the vulnerabilities Rust cannot prevent. Within this paper 68 prominent open source projects, as linux kernel and Curl, were analyzed in order to map the vulnerabilities.

The study identified a total of 4,507 vulnerabilities across these projects, 2,257 of which had specific Common Weakness Enumeration CWE designations. These were mapped to Rust's safety guarantees to determine preventative potential. The analysis reveals that a significant majority of historical security issues in C are memory safety errors that Rust effectively eliminates at compile time.




Figure 2.2: Rust Preventions based on CWE (Meneely et al. 2025)

As illustrated in Figure 2.2, approximately 56.9% of these historical vulnerabilities fall into the "Virtually Impossible" category, meaning they would be prevented by Rust's strict type system and borrow checker without requiring runtime overhead. Conversely, 29.3% of vulnerabilities—primarily logic errors or domain-specific flaws—remain unpreventable by the language choice alone (Meneely et al. 2025).

Table 1 details the top 20 most common vulnerability types found in the subject projects. The data confirms that the two most frequent issues, NULL Pointer Dereference (CWE-476) and Use After Free (CWE-416), accounting for nearly 17% of all analyzed CVEs, are completely preventable in safe Rust (Meneely et al. 2025).

CWE ID	Vulnerability Name	# CVEs	# Proj	Rust Prevention
CWE-476	NULL Pointer Dereference	411	17	Virtually Impossible
CWE-416	Use After Free	343	17	Virtually Impossible
CWE-125	Out-of-bounds Read	145	14	Virtually Impossible
CWE-20	Improper Input Validation	105	17	No Help
CWE-190	Integer Overflow or Wraparound	100	12	Discouraged (Debug)
CWE-401	Memory Leak	96	6	Virtually Impossible
CWE-787	Out-of-bounds Write	92	15	Virtually Impossible
CWE-667	Improper Locking	86	4	Opt-In Measures
CWE-369	Divide By Zero	78	4	Virtually Impossible
CWE-617	Reachable Assertion	63	6	Discouraged (Debug)
CWE-362	Race Condition	62	11	Discouraged (Design)
CWE-119	Improper Buffer Restriction	62	14	Virtually Impossible
CWE-400	Uncontrolled Resource Consumption	59	15	No Help
CWE-122	Heap-based Buffer Overflow	53	12	Virtually Impossible
CWE-200	Exposure of Sensitive Information	45	11	No Help
CWE-120	Classic Buffer Overflow	45	9	Virtually Impossible
CWE-908	Use of Uninitialized Resource	35	5	Virtually Impossible
CWE-415	Double Free	34	8	Virtually Impossible
CWE-770	Resource Allocation without Limits	29	7	No Help
CWE-131	Incorrect Buffer Size Calculation	26	7	Virtually Impossible

Table 2.2: Top 20 Most Common Vulnerabilities and Rust Prevention Status
(Meneely et al. 2025)

This data suggests that while Rust is not a "silver bullet" for all security issues (such as input validation or logic bugs), its adoption could have theoretically prevented the majority of the most severe historical vulnerabilities found in these critical infrastructure projects (Meneely et al. 2025).

Chapter 3

Methodology and Solution Proposal

The migration of critical software infrastructure from C++ to Rust is not a syntactic translation exercise but a fundamental architectural transformation. It requires rethinking ownership, concurrency, and interface boundaries in order to preserve functionality while improving safety guarantees.

As discussed in Chapter 2, the so-called “Boost Paradox” introduces a unique challenge: the very features that make Boost powerful—such as templates and metaprogramming—also hinder the applicability of automated interoperability tools.

This chapter defines a structured methodology to decompose Boost-based systems and proposes an architectural design that incrementally bridges the safety gap between C++ and Rust.

3.1 Component Selection Criteria

Historically, a “big bang” rewrite of large C++ codebases into Rust has proven prone to failure due to the difficulty of verifying feature parity, the disruption of ongoing development, and the elevated certification risk in safety-critical domains. Consequently, this research adopts a selective and incremental migration strategy.

Boost libraries are not considered uniformly; instead, they are evaluated and selected based on their relevance, isolation capability, and suitability for automotive software architectures.

3.1.1 Automotive-Oriented Selection Rationale

The automotive sector has increasingly adopted open-source and modular software platforms to address growing system complexity and multi-vendor integration. In this context, several automotive manufacturers and suppliers participate in Eclipse Software Defined Vehicle initiatives, including Eclipse Safety-Critical Open Research Environment (SCORE).

Eclipse SCORE targets deterministic, portable, and interoperable software architectures with long-term support guarantees—key concerns for ISO-26262 compliant systems (Foundation 2025). These initiatives reflect the realities of modern automotive development: large, heterogeneous C++ codebases, long product lifecycles, and strict constraints on timing, memory

usage, and failure modes. Such characteristics render wholesale language replacement infeasible and reinforce the need for incremental, well-contained migration strategies.

The Boost libraries selected for migration are therefore evaluated according to the following automotive-oriented criteria:

- **Relevance to in-vehicle software:** Preference is given to libraries commonly used in middleware, communication, configuration management, and concurrency control, which are prevalent in automotive Electronic Control Units (ECUs) and service-oriented vehicle architectures.
- **Deterministic behavior:** Libraries whose runtime behavior is predictable and suitable for real-time or near-real-time constraints are prioritized, aligning with automotive requirements for bounded latency and reproducible execution.
- **Safety and robustness:** Components that encode disciplined resource management or concurrency patterns (e.g., `Boost.SmartPtr`, `Boost.Thread`, `Boost.Asio`) are preferred, as they conceptually align with Rust's ownership and concurrency models.
- **Isolation capability:** Libraries that can be encapsulated behind stable, well-defined interfaces are favored, allowing them to serve as clear boundaries for Foreign Function Interface (FFI) integration.
- **Industrial adoption and tooling maturity:** Libraries actively used or referenced in automotive-oriented open-source projects, such as those under the Eclipse SCORE ecosystem, are considered more suitable due to their demonstrated industrial relevance.

3.1.2 Deprioritized Components

Not all Boost libraries are suitable for migration or interoperability with Rust. Libraries that rely heavily on advanced template metaprogramming, domain-specific embedded languages, or extensive compile-time code generation are deprioritized due to their limited analyzability and poor suitability for FFI-based integration.

Similarly, libraries not directly used or referenced in automotive-oriented open-source projects such as Eclipse SCORE are not considered immediate migration candidates within the scope of this work.

3.1.3 Safety-Critical Constraints and ISO 26262 Alignment

ISO-26262 requires that software architectures for automotive safety systems support freedom from interference, predictable execution, and systematic fault avoidance. Programming languages and libraries used in such systems must minimize undefined behavior, provide explicit ownership semantics, and enable static verification wherever possible (ISO 2018).

C++'s permissive memory model and reliance on developer discipline pose challenges when targeting higher Automotive Safety Integrity Levels (ASIL) levels. In practice, Boost libraries

are frequently employed in automotive systems to impose structure and consistency on C++ development, particularly in resource management, concurrency, and asynchronous communication (Synopsys 2025).

The proposed methodology leverages this disciplined use of Boost as a preparatory step toward Rust integration. By constraining C++ systems to well-defined Boost abstractions, reliance on raw pointers, implicit ownership conventions, and ad hoc concurrency mechanisms is reduced. This aligns with ISO-26262 recommendations for limiting systematic faults through controlled language subsets, coding guidelines, and architectural enforcement (ISO 2018).

Rust components introduced under this methodology are designed to be ownership-complete and memory-safe by construction. Unsafe code is strictly isolated at Foreign Function Interface (FFI) boundaries, supporting safety argumentation, auditability, and certification activities.

3.1.4 Boost Libraries Relevant to Eclipse SCORE

In Eclipse SCORE-style automotive architectures, Boost libraries are commonly used to support middleware services, configuration management, concurrency, and asynchronous I/O. Their selection is driven by the need to enforce structure and predictability in large-scale, safety-relevant C++ systems.

Table 3.1 summarizes Boost libraries commonly applicable to automotive platforms such as Eclipse SCORE and their corresponding Rust-native equivalents.

This mapping highlights that several Boost libraries already encode concepts such as explicit ownership, variant-based state, and structured concurrency—concepts that are native to Rust. As a result, the primary challenge of migration is architectural rather than conceptual, reinforcing the feasibility of incremental Rust adoption in safety-critical automotive ecosystems.

3.1.5 Comparative Analysis as a Migration Enabler

A central step in the proposed methodology is the comparison between Boost libraries used in automotive systems and their Rust-native equivalents. This comparison is not a simple replacement exercise, but an architectural analysis that exposes semantic gaps, safety assumptions, and ecosystem limitations.

By explicitly mapping Boost abstractions to Rust constructs, developers can:

- Identify which safety properties are already enforced at the type level
- Detect mismatches in ownership, error handling, or concurrency semantics
- Quantify the residual risk that remains after migration

Boost Library	Rust Equivalent	Automotive Relevance
Boost.Asio	tokio, async-std	Asynchronous I/O and event handling for in-vehicle services and middleware; requires careful runtime integration in safety-critical contexts
Boost.Thread	std::thread, crossbeam	Thread management and synchronization for concurrent ECU software components
Boost.SmartPtr	Box, Rc, Arc	Explicit ownership and lifetime management; directly maps to Rust's ownership model
Boost.Optional	Option<T>	Explicit representation of optional values, reducing NULL pointer usage
Boost.Variant	enum	Type-safe sum types used in configuration, messaging, and state machines
Boost.Program_options	clap, argh	Configuration parsing for tooling and non-safety-critical control components
Boost.Filesystem	std::fs, std::path	File and path handling in diagnostic, logging, and update subsystems
Boost.Chrono	std::time	Time measurement and timeout handling with deterministic semantics

Table 3.1: Boost libraries relevant to Eclipse SCORE-style automotive systems and Rust equivalents

This analysis also reveals areas where no suitable Rust equivalent exists for automotive-grade use. In such cases, the absence of a mature Rust library becomes an explicit design signal rather than an implementation obstacle.

3.1.6 From Library Gaps to New Rust Automotive Crates

The comparison between Boost and Rust ecosystems frequently uncovers functionality that is either missing or insufficiently mature in Rust for safety-critical automotive contexts.

Rather than forcing unsafe bindings or retaining C++ dependencies indefinitely, the proposed methodology treats these gaps as opportunities for new Rust library development. Such libraries are designed from the outset with:

- Explicit ownership and lifetime semantics
- Deterministic execution models suitable for real-time analysis
- Minimal and auditable unsafe code
- Alignment with AUTOSAR Adaptive concepts

In this sense, Boost serves not only as a migration bridge, but also as a functional benchmark that informs the evolution of Rust's automotive ecosystem. The migration process thus becomes bidirectional: existing C++ designs guide Rust library development, while Rust's guarantees inform safer architectural patterns in C++.

3.1.7 ASIL-Oriented Migration Strategy

ISO-26262 distinguishes software components according to their ASIL, ranging from Quality Management (QM) to ASIL-D. Each level imposes increasing requirements on fault avoidance, fault detection, and freedom from interference. As a result, a uniform migration strategy is neither practical nor desirable across all ASIL levels.

The proposed methodology adopts a differentiated migration approach:

- **QM and ASIL-A components:** Tooling, diagnostics, logging, configuration services, and non-safety-critical middleware. These components are suitable early candidates for Rust adoption, as certification constraints are less restrictive.
- **ASIL-B and ASIL-C components:** Communication services, resource managers, and coordination logic. Rust integration at this level focuses on eliminating memory safety defects while maintaining deterministic execution. Unsafe Rust is strictly isolated and subject to additional verification.
- **ASIL-D components:** Rust is introduced only where its guarantees provide demonstrable safety benefits. Components must be ownership-complete, free from dynamic allocation where required, and analyzable under Worst Case Execution Time (WCET) constraints. FFI boundaries are minimized and treated as safety-relevant interfaces subject to rigorous review.

This stratification aligns with ISO-26262 recommendations for mixed-criticality systems and supports freedom from interference between software elements of different ASIL levels.

3.1.8 Migration Workflow Overview

Figure 3.1 presents the proposed incremental migration workflow for integrating Rust into Boost-based automotive software systems. The workflow is designed to support safety-critical development under ISO-26262 constraints while preserving system continuity and certification traceability.

The process begins with a **comparative analysis** of existing Boost library usage and available Rust-native equivalents. This step establishes a semantic baseline by identifying ownership models, concurrency assumptions, error-handling mechanisms, and determinism properties already enforced in the C++ architecture.

Based on this analysis, an **architectural refactoring phase** follows, during which Boost abstractions are used to constrain C++ components into well-defined, analyzable interfaces. This

step reduces semantic ambiguity, limits undefined behavior, and prepares the system for language boundary introduction without altering functional behavior.

Next, **FFI boundary definition** is performed. Stable and minimal interfaces are identified between C++ and Rust components, with explicit ownership transfer rules and data lifetime guarantees. These boundaries are treated as safety-relevant interfaces and are subject to additional review and verification.

The workflow then proceeds with **incremental Rust integration**, where Rust components are introduced selectively according to their assigned ASIL level. Initial integration targets QM and ASIL-A components, followed by progressively higher-criticality elements as confidence, tooling maturity, and safety argumentation evolve.

Throughout the process, **continuous verification and validation** are performed to ensure functional equivalence, timing predictability, and freedom from interference. Feedback from each iteration informs subsequent migration steps, allowing the workflow to adapt to architectural constraints and certification requirements.




Figure 3.1: Incremental migration workflow for Boost-based automotive systems integrating Rust under ISO-26262 constraints

Overall, the workflow emphasizes early reduction of semantic ambiguity on the C++ side and controlled introduction of Rust components, enabling safer and more predictable interoperability while maintaining compliance with automotive safety standards.

Chapter 4

Development Plan

This chapter presents the development plan for the proposed Boost-to-Rust migration methodology. It details the planned reimplementation of selected components in Rust, the design of stable interfaces, and the integration of bidirectional FFI. The plan is structured to support incremental adoption while maintaining compliance with automotive safety standards such as ISO-26262 and architectural principles from AUTOSAR Adaptive.

The development activities are scheduled between February and July and are organized to progressively reduce technical risk while preserving functional continuity of the existing C++ system.

4.1 Objectives and Scope

The primary objectives of this development phase are:

- To reimplement selected Boost-based components in Rust with equivalent functionality and improved memory-safety guarantees.
- To design explicit, minimal, and auditable interfaces between C++ and Rust components.
- To integrate Rust modules incrementally using bidirectional FFI, preserving deterministic behavior and freedom from interference.
- To validate the resulting hybrid system against safety, correctness, and performance expectations relevant to automotive software.

The scope of this chapter is limited to architectural and integration-level development. Functional feature expansion beyond parity with the existing C++ implementation is explicitly excluded.

4.2 Development Phases Overview

The development plan is divided into six sequential phases:

1. System analysis and Boost component selection
2. Architectural refactoring and interface definition
3. Rust reimplementation of selected components
4. Bidirectional FFI integration
5. Validation and safety analysis
6. Consolidation and documentation

Each phase produces concrete artifacts that serve as inputs for the subsequent phase, enabling traceability and controlled progression consistent with ISO-26262 development practices.

4.3 System Analysis and Component Selection

February

The initial phase focuses on analyzing the existing C++ system and identifying Boost libraries that are relevant to automotive software architectures such as those used in Eclipse SCORE-based environments.

The analysis includes:

- Identification of Boost libraries in use and their functional roles
- Classification of components by criticality and ASIL level
- Mapping of Boost abstractions to potential Rust-native equivalents

This phase establishes the baseline architecture and defines the migration targets for subsequent development.

4.4 Architectural Refactoring and Interface Design

March

In this phase, the selected C++ components are refactored to reduce semantic ambiguity and prepare them for interoperability. Boost abstractions are used to replace raw pointers, ad hoc ownership conventions, and implicit lifetimes.

Key activities include:

- Encapsulation of migrated components behind stable interfaces
- Elimination of undefined behavior patterns
- Definition of ownership and lifetime semantics at component boundaries

This refactoring step does not change observable behavior but improves analyzability and safety argumentation.

4.5 Rust Reimplementation

April

During this phase, selected Boost-based components are reimplemented in Rust. The Rust implementations are designed to be ownership-complete and memory-safe by construction.

Design constraints include:

- Explicit ownership and borrowing semantics
- Deterministic execution suitable for timing analysis
- Minimal use of unsafe code, isolated and documented

Functional equivalence with the original C++ components is verified through unit testing and interface-level validation.

4.6 Bidirectional FFI Integration

May

This phase introduces bidirectional FFI integration between C++ and Rust. Stable C-compatible interfaces are defined, and responsibility for memory allocation, ownership transfer, and error handling is made explicit.

Key concerns addressed include:

- Isolation of unsafe code at FFI boundaries
- Prevention of exception and panic propagation across language borders
- Alignment with AUTOSAR Adaptive service-oriented communication models

The FFI boundary is treated as a safety-relevant interface and is subject to additional review.

4.7 Validation and Safety Analysis

June

Validation activities focus on ensuring that the hybrid C++/Rust system meets functional, safety, and performance expectations.

This includes:

- Functional regression testing

- Concurrency and race-condition analysis
- Review of freedom from interference between components of different ASIL levels

The results of this phase support the safety argumentation required for ISO-26262 compliant systems.

4.8 Consolidation and Documentation

July

The final phase consolidates the development results and prepares the system for evaluation and reporting.

Activities include:

- Final architectural documentation
- Evaluation of migration feasibility and limitations
- Documentation of lessons learned and future work

This phase ensures that the development process and outcomes are reproducible and auditable.

4.9 Development Timeline

Table 4.1 presents a Gantt-style overview of the development plan covering the February–July period.

Activity / Month	Feb	Mar	Apr	May	Jun	Jul
System Analysis and Boost Selection	X					
Architectural Refactoring		X				
Rust Reimplementation			X			
Bidirectional FFI Integration				X		
Validation and Safety Analysis					X	
Consolidation and Documentation						X

Table 4.1: Gantt-style development plan for the Boost-to-Rust migration

Bibliography

70% of Security Bugs Are Memory Safety Issues (2020). Tech. rep. Microsoft Security Response Center.

Amazon Web Services (2026). *Firecracker: Secure and Fast MicroVMs for Serverless Computing*. Firecracker main page. url: <https://firecracker-microvm.github.io/>.

AUTOSAR Partnership (2023). *AUTOSAR Adaptive Platform - Explanation of Adaptive Platform Design*. Tech. rep. R23-11. Describes the service-oriented C++ architecture used in modern SDVs. AUTOSAR GbR. url: <https://www.autosar.org/standards/adaptive-platform>.

Blandy, Jim, Jason Orendorff, and Leonora F.S. Tindall (2021). *Programming Rust: Fast, Safe Systems Development*. Sebastopol, CA: O'Reilly Media. isbn: 978-1-492-05259-3. url: <https://oreilly.com>.

Boost C++ Libraries Documentation (2025). url: <https://www.boost.org/docs/>.

Bristol, Andrea (2025). *Ada and Rust are highlighted by the NSA and CISA in Memory Safe Language Information Sheet*. url: <https://www.adacore.com/blog/ada-and-rust-are-highlighted-by-the-nsa-and-cisa-in-memory-safe-language-information-sheet>.

C++ ABI Compatibility (2025). Accessed: 2025-01-18. url: https://parallel-rust-cpp.github.io/cpp_abi.html.

CppCon (2020). *CppCon 2020 Presentation Materials*. Cpp Conference. url: <https://github.com/CppCon/CppCon2020>.

crates.io: The Rust Community's Crate Registry (2025). url: <https://crates.io/crates?sort=recent-downloads>.

Dashdamirli, N. (2025). *Analyzing the Performance and Practicality of C, Rust, Python, and Lua Programming Languages for Developing Microcontroller Applications*. Technical report.

Deloitte (2024). *Software-defined vehicles: Global manufacturer readiness study*. url: <https://www.deloitte.com/ua/en/Industries/automotive/analysis/software-defined-vehicles.html>.

Foundation, Eclipse (2025). *Eclipse Safe Open Vehicle Core (S-Core)*. url: <https://projects.eclipse.org/projects/automotive.score>.

Google (2021). *High-level design of C++/Rust interop*. States impossibility of using C++ templates via standard bindgen/C FFI. url: <https://crubit.rs/design/design.html>.

Google Android Team (2024). *Rust / C++ Interop in Android Platform*. Case study of large-scale hybrid C++/Rust integration in a production OS. url: <https://source.android.com/docs/setup/build/rust/building-rust-modules/cpp-interop>.

Hong, J. (2023). "Improving Automatic C-to-Rust Translation with Static Analysis". In: *2023 IEEE/ACM 45th International Conference on Software Engineering: Companion Proceedings (ICSE-Companion)*.

- Immunant, Inc. (2025). *c2rust: Migrate C code to Rust*. Version 0.21.0. url: <https://github.com/immunant/c2rust>.
- ISO (2018). *ISO 26262:2018 Road vehicles – Functional safety*. Part 6: Product development at the software level. Geneva, Switzerland: International Organization for Standardization.
- ISO 26262-6:2018 Road vehicles — Functional safety — Part 6: Product development at the software level* (2018). Defines requirements for software unit testing and freedom from interference. Geneva, Switzerland: International Organization for Standardization. url: <https://www.iso.org/standard/68388.html>.
- ISO/IEC 14882:2020 Programming languages — C++* (Dec. 2020). Defines the C++20 standard, including modules, coroutines, concepts, and ranges. Geneva, Switzerland: International Organization for Standardization. url: <https://www.iso.org/standard/79358.html>.
- Khan, Mustafif (2023). *Rust for C++ Programmers: Learn how to embed Rust in C/C++ with ease*. London, UK: BPB Online. isbn: 978-93-5551-355-7. url: <https://www.bpbonline.com>.
- Klabnik, Steve, Carol Nichols, and The Rust Project Developers (2017). *The Rust Programming Language*. San Francisco, CA: No Starch Press. isbn: 978-1593278281. url: <https://doc.rust-lang.org/book/>.
- Kuo, Akhil (2024). *Rust for Blockchain Application Development: Use the Power of Rust to Develop Scalable, Robust, and Secure Blockchain Applications*. Birmingham, UK: Packt Publishing. isbn: 978-1-83763-464-4. url: <https://www.packtpub.com/en-us/product/rust-for-blockchain-application-development-9781837634644>.
- Ling, M. et al. (2022). "In Rust We Trust: A Transpiler from Unsafe C to Safer Rust". In: *2022 IEEE/ACM 44th International Conference on Software Engineering: Companion Proceedings (ICSE-Companion)*. Pittsburgh, PA, USA, pp. 354–355.
- Memory Safety in Android and Chrome* (2022). Tech. rep. Google Project Zero.
- Meneely, A. et al. (2025). "Just Use Rust: A Best-Case Historical Study of Open Source Vulnerabilities in C". In: *2025 IEEE/ACM 3rd International Workshop on Software Vulnerability Management (SVM)*.
- Mozilla (2026). *Including Rust Code in Firefox*. Firefox Source Docs documentation. Mozilla. url: <https://firefox-source-docs.mozilla.org/build/buildsystem/rust.html>.
- Okawa, S. and S. Yamaguchi (2024). "A Performance Study on Rust and C Programs". In: *2024 Twelfth International Symposium on Computing and Networking (CANDAR)*.
- Rooney, M. P. and S. J. Matthews (2023). "Evaluating FFT Performance of the C and Rust Languages on Raspberry Pi Platforms". In: *2023 IEEE International Conference on Cluster Computing (CLUSTER)*.
- Rust Documentation* (2025). Accessed: 2025-01-18. url: <https://docs.rs/>.
- S&P Global Mobility (2025). *The software-defined vehicle market is taking off*. url: <https://www.spglobal.com/automotive-insights/en/blogs/2025/09/the-software-defined-market-is-taking-off>.
- Stroustrup, Bjarne (2013). *The C++ Programming Language*. 4th ed. Addison-Wesley.

- (2020). “Thriving in a Crowded and Changing World: C++ 2006–2020”. In: *Proceedings of the ACM on Programming Languages* 4.HOPL, 70:1–70:168. doi: 10.1145/3386320. url: <https://www.stroustrup.com/hopl20main-p5-p-bfc9cd4--final.pdf>.
- Synopsys (2025). *What is ASIL (Automotive Safety Integrity Level)?* Accessed: 2026-01-31. url: <https://www.synopsys.com/glossary/what-is-asil.html>.
- The Linux Kernel Organization (2026). *Rust — The Linux Kernel documentation*. Documentation for experimental Rust support in the Linux kernel (v6.1+). url: <https://docs.kernel.org/rust/index.html>.
- The MISRA Consortium (Oct. 2023). *MISRA C++:2023 Guidelines for the use of C++17 in critical systems*. Industry standard attempting to mitigate C++ memory safety risks via subsetting. MISRA Ltd. Nuneaton, UK. isbn: 978-1-906400-33-3. url: <https://misra.org.uk/>.
- The Rust Project Developers (2026). *Embedded Devices - Rust Programming Language*. Rust for embedded devices. url: <https://www.rust-lang.org/what/embedded>.
- The rust-bindgen Project Developers (2026). *Crate bindgen: Generate Rust bindings for C and C++ libraries*. Rust documentation. url: <https://docs.rs/bindgen/latest/bindgen/>.
- Tolnay, David (2025). *cxx: Safe Interop between Rust and C++*. The industry-standard library for safe, static-analysis-driven FFI between Rust and C++. url: <https://github.com/dtolnay/cxx>.
- Xia, L., Y. Wu, and B. Hua (2023). “Rustcheck: Safety Enhancement of Unsafe Rust via Dynamic Program Analysis”. In: *2023 IEEE 23rd International Conference on Software Quality, Reliability, and Security Companion (QRS-C)*.
- Zakorum (2023). *boosters: Rust port of boost libraries for C++*. Version 0.1.0. Rust crate version 0.1.0. url: <https://crates.io/crates/boosters>.